

Installing Your First Big Data Cluster Using Cloudera CDH



Xavier Morera

PASSIONATE ABOUT TEACHING

@xmorera www.xaviermorera.com



The Three Cloudera Installation Paths

A large, bold, blue capital letter 'A' representing Path A.

Path A

A large, bold, maroon capital letter 'B' representing Path B.

Path C

A large, bold, green capital letter 'C' representing Path C.

Path C



The Three Cloudera Installation Paths

Non Production



Path A



Path B

Production



Path B



Path C



Installation Path A



Automated installation

Using Cloudera Manager

- Cloudera Manager Server
- JDK, Agents, CDH and services

Automatic embedded PostgreSQL

Mainly for POCs, dev and demos

Not intended for production

- Scaling limitations
- May require database migration



Installation Path B

B

Automated or manual installation

Parcels or Packages

Puppet or Chef to help install across cluster

Production and Non-Production

Non-Production

- Embedded PostgreSQL

Production

- Manual database installation & configuration



Installation Path C



Manual installation

- JDK, Cloudera Manager Server & Agents
- Tarballs used for installation

Use Cloudera Manager to install and deploy

- CDH
- Managed software services (parcels)



Next Step?



Let's Pick One
Installation Path A




```
wget https://archive.cloudera.com/  
cm5/installer/latest/cloudera-manager-installer.bin
```

Getting the Installation File

Latest version

Find the one that you need

1



```
chmod u+x cloudera-manager-installer.bin
```

Execute Permissions

chmod gives execute permissions to the user

.bin is a self extracting binary file

Execute this file to install **Cloudera Manager**



```
sudo ./cloudera-manager-installer.bin
```

Install Cloudera Manager

Execute **.bin** file which will start the process

Notice the **./** and **sudo** is required

Use **--skip_repo_package=1** for local package installation

3




Demo



Exercise 7

Installing Cloudera Manager using Path A





- cloudera-quickstart-v...**
Saved
- Base Linux Machine (F...)**
Powered Off
- mgt01 (mgt01 clean)**
Powered Off
- dn01 (clean dn01)**
Powered Off
- dn02 (clean dn02)**
Powered Off
- dn03 (clean dn03)**
Powered Off
- dn04 (clean dn04)**
Powered Off

General

Name: mgt01
Operating System: Red Hat (64-bit)

System

Base Memory: 2048 MB
Processors: 2
Boot Order: Floppy, Optical, Hard Disk
Acceleration: VT-x/AMD-V, Nested Paging, PAE/NX, KVM Paravirtualization

Display

Video Memory: 16 MB
Remote Desktop Server: Disabled
Video Capture: Disabled

Storage

Controller: IDE
IDE Primary Master: [Optical Drive] Empty
IDE Secondary Master: [Optical Drive] Empty
Controller: SATA
SATA Port 0: mgt01.vdi (Normal, 80.00 GB)

Audio

Host Driver: Windows DirectSound
Controller: ICH AC97

Network

Adapter 1: Intel PRO/1000 MT Desktop (NAT Network, 'PSNATNetwork')

USB

Shared folders


None

Description

Preview

mgt01

Installation Step Two: Deploying Your Cluster



Cloudera Manager is installed

Deploy agents to all nodes in our cluster

Automated installation

- But you need to make a few decision

Cloudera Manager provides

- Easy to understand interface
- Detailed log during installation

What's left?



Steps



Select Cloudera Manager version

Specify hosts

Select repository

Accept licenses for JDK installation

Single user mode

Login credentials

Simultaneous installations



Steps



Parcels: downloaded, distributed, unpacked and activated

Inspect hosts for correctness

Select combination of services to install

Select nodes for services

Database setup

Configuration parameters

Cluster setup complete!



Demo



Exercise 8

Deploying Big Data Cluster Using Cloudera CDH Automated Install (Path A)



☐ Remember me

Takeaway



Set up a Hadoop cluster

- 5 nodes

CDH

- Cloudera Distribution including Hadoop

Cloudera Manager

Automated install

- Path A
- Multiple steps
- Embedded database
- Non prod but works great for dev

Great first step to Big Data

